

# 网络空间安全实训实验报告

纪盛谦 57118218 2021-7-22

## Task1 : Implementing a Simple Firewall

### Task1.A : Implement a Simple Kernel Module

将其 make 后，加载、移除 hello 后，可以看到有记录。

```
[07/22/21]seed@VM:~/.../kernel_module$ sudo insmod hello.ko
[07/22/21]seed@VM:~/.../kernel_module$ lsmod | grep hello
hello                16384  0
[07/22/21]seed@VM:~/.../kernel_module$ sudo rmmod hello
```

```
[ 379.801387] Hello World!
[ 401.623965] Bye-bye World!.
```

### Task 1.B: Implement a Simple Firewall Using Netfilter

(1) 挂载程序后，到 8.8.8.8/53 的报文被拦截，移除检测程序后则可以发出该报文且有回复。

```
[07/22/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
[07/22/21]seed@VM:~/.../packet_filter$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; connection timed out; no servers could be reached

[07/22/21]seed@VM:~/.../packet_filter$ sudo rmmod seedFilter
[07/22/21]seed@VM:~/.../packet_filter$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 24954
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags;; udp: 512
;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                20721   IN      A      93.184.216.34

;; Query time: 56 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
```

---

```

[ 2532.012294] 10.0.2.15 --> 8.8.8.8 (UDP)
[ 2532.012299] *** Dropping 8.8.8.8 (UDP), port 53
[ 2537.009091] *** LOCAL_OUT
[ 2537.009093] 10.0.2.15 --> 8.8.8.8 (UDP)
[ 2537.009103] *** Dropping 8.8.8.8 (UDP), port 53
[ 2542.005465] *** LOCAL_OUT
[ 2542.005467] 10.0.2.15 --> 8.8.8.8 (UDP)
[ 2542.005476] *** Dropping 8.8.8.8 (UDP), port 53
[ 2572.557543] The filters are being removed.

```

(2) 通过实验可以发现数据报从进入系统，进行 IP 校验以后，首先经过第一个 HOOK 函数 `NF_IP_PRE_ROUTING` 进行处理；然后就进入路由代码，其决定该数据报是需要转发还是发给本机的；若该数据报是发给本机的，则该数据报经过 HOOK 函数 `NF_IP_LOCAL_IN` 处理以后然后传递给上层协议；若该数据报应该被转发则它被 `NF_IP_FORWARD` 处理；经过转发的数据报经过最后一个 HOOK 函数 `NF_IP_POST_ROUTING` 处理以后，再传输到网络上。

本地产生的数据经过 HOOK 函数 `NF_IP_LOCAL_OUT` 处理后，进行路由选择处理，然后经过 `NF_IP_POST_ROUTING` 处理后发送出去。

```

[ 130.608683] Registering filters.
[ 133.777310] *** LOCAL_OUT
[ 133.777312] 10.0.2.15 --> 8.8.8.8 (UDP)
[ 133.777320] *** POST_ROUTING
[ 133.777320] 10.0.2.15 --> 8.8.8.8 (UDP)
[ 133.825915] *** PRE_ROUTING
[ 133.825928] 8.8.8.8 --> 10.0.2.15 (UDP)
[ 133.825940] *** LOCAL_IN
[ 133.825944] 8.8.8.8 --> 10.0.2.15 (UDP)
[ 133.826875] *** LOCAL_OUT
[ 133.826876] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 133.826880] *** POST_ROUTING
[ 133.826881] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 133.826886] *** PRE_ROUTING
[ 133.826887] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 133.826888] *** LOCAL_IN
[ 133.826888] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 133.826973] *** LOCAL_OUT
[ 133.826973] 10.0.2.15 --> 8.8.8.8 (UDP)
[ 133.826976] *** POST_ROUTING
[ 133.826977] 10.0.2.15 --> 8.8.8.8 (UDP)

```

(3) 选用 `NF_IP_LOCAL_IN` 这个 hook，最终实验结果如下，从 10.9.0.5 向 10.9.0.1 ping 和 telnet 都不能通。

```

root@0a912ff6278e:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^Z
[1]+  Stopped                  ping 10.9.0.1
root@0a912ff6278e:/# telnet 10.9.0.1
Trying 10.9.0.1...
^Z^C
root@0a912ff6278e:/# █

```

其中处理函数代码如下

```
unsigned int block(void *priv, struct sk_buff *skb,
                  const struct nf_hook_state *state)
{
    struct iphdr *iph;
    struct tcphdr *tcph;

    u16 port = 23;

    if (!skb) return NF_ACCEPT;

    iph = ip_hdr(skb);

    if (iph->protocol == IPPROTO_ICMP)
    {
        return NF_DROP;
    }

    if (iph->protocol == IPPROTO_TCP)
    {
        tcph = tcp_hdr(skb);
        if (ntohs(tcph->dest) == port)
        {
            return NF_DROP;
        }
    }

    return NF_ACCEPT;
}
```

## Task2: Experimenting with Stateless Firewall Rules

### Task2.A: Protecting the Router

在 router 上设置 iptable，如果只安装手册的指令是不能达到效果的，只有有在 INPUT 和 OUTPUT 都设置通过 echo-request 和 echo-reply 才可以。

```
root@f71f0b983d48:/# iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@f71f0b983d48:/# iptables -A OUTPUT -p icmp --icmp-type echo-request -j ACCEPT
root@f71f0b983d48:/# iptables -P OUTPUT DROP
root@f71f0b983d48:/# iptables -P INPUT DROP
root@f71f0b983d48:/# iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@f71f0b983d48:/# iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
```

最终从 10.9.0.5 向路由 ping 和 telnet 效果如下。

```
root@0a912ff6278e:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.049 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.049 ms
^Z
[4]+  Stopped                  ping 10.9.0.11
root@0a912ff6278e:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
```

## Task2.B: Protecting the Internal Network

在路由上配置 iptable

```
root@f71f0b983d48:/# iptables -A FORWARD -o eth1 -p icmp --icmp-type echo-request -j DROP
root@f71f0b983d48:/# iptables -A FORWARD -i eth1 -p icmp --icmp-type echo-reply -j DROP
root@f71f0b983d48:/# iptables -A FORWARD -i eth0 -p icmp --icmp-type echo-request -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -o eth0 -p icmp --icmp-type echo-reply -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -o eth1 -p icmp --icmp-type echo-reply -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -i eth1 -p icmp --icmp-type echo-request -j ACCEPT
root@f71f0b983d48:/# iptables -P FORWARD DROP
```

外部主机无法 ping 内部主机

```
root@0a912ff6278e:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
^Z
[1]+  Stopped                  ping 192.168.60.5
```

外部主机可以 ping 路由

```
root@0a912ff6278e:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.047 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.069 ms
^Z
[2]+  Stopped                  ping 10.9.0.11
```

内部主机可以 ping 外部主机

```
root@0a912ff6278e:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.047 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.069 ms
^Z
[2]+  Stopped                  ping 10.9.0.11
```

其他数据包内外不通（内外无法互相 telnet）

```
root@0a912ff6278e:/# telnet 192.168.60.5
Trying 192.168.60.5...
^C
-

root@0a912ff6278e:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
-
```



## Task 2.C: Protecting Internal Servers

在路由上配置 iptable

```
root@f71f0b983d48:/# iptables -A FORWARD -i eth0 -p tcp --dport 23 -d 192.168.60.5 -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -o eth0 -p tcp --sport 23 -s 192.168.60.5 -j ACCEPT
root@f71f0b983d48:/# iptables -P FORWARD DROP
```

外部主机可以 telnet 到 192.168.60.5

```
root@0a912ff6278e:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
660e8067c2c5 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

外部主机无法 telnet 到其他内部主机

```
root@0a912ff6278e:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
^C
```

内部主机可以 telnet 到其他内部主机

```
root@660e8067c2c5:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
60979ddce039 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

内部主机无法 telnet 到外部主机

```
[07/22/21]seed@VM:~$ docksh 66
root@660e8067c2c5:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
```

## Task3: Connection Tracking and Stateful Firewall

### Task 3.A: Experiment with the Connection Tracking

ICMP 连接大概持续 30s

```
icmp      1 28 src=192.168.60.1 dst=192.168.60.5 type=8 code=0 id=4 src=192.168.60.5 dst=192.168.60.1 type=0 code=0 id=4 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 3 flow entries have been shown.
```

UDP 连接也大概持续 30s

```
udp       17 28 src=192.168.60.1 dst=192.168.60.5 sport=37432 dport=9090
ED] src=192.168.60.5 dst=192.168.60.1 sport=9090 dport=37432 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 2 flow entries have been shown.
```

TCP 连接大概持续 120s

```
tcp       6 116 SYN_SENT src=192.168.60.1 dst=192.168.60.5 sport=34338 dport=34338
[UNREPLIED] src=192.168.60.5 dst=192.168.60.1 sport=9090 dport=34338 mark=1
```

### Task 3.B: Setting Up a Stateful Firewall

在路由上配置 iptable

```
root@f71f0b983d48:/# iptables -A FORWARD -p tcp -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -p tcp -i eth0 -d 192.168.60.5 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -p tcp -o eth1 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -p tcp -o eth0 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
root@f71f0b983d48:/# iptables -A FORWARD -p tcp -i eth1 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
root@f71f0b983d48:/# iptables -P FORWARD DROP
```

外部主机可以 telnet 到 192.168.60.5

```
root@0a912ff6278e:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
660e8067c2c5 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

外部主机无法 telnet 到其他内部主机

```
root@0a912ff6278e:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
^C
```

内部主机可以 telnet 到其他内部主机

```
root@660e8067c2c5:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
60979ddce039 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

内部主机可以 telnet 到外部主机

```
root@660e8067c2c5:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
0a912ff6278e login: seed
```

二者优缺点：基于连接的防火墙只需要在建立连接的时候判定是否合法，之后的报文只需要判定是否建立连接即可，但它需要调用 `conntrack`；而没有基于连接的防火墙则需要对所有报文进行判定是否合法，但不需要借助 `conntrack`。

## Task4: Limiting Network Traffic

两条都在的情况下，一开始会有 5 个比较快的 ping，之后平均每分钟 ping10 个报文。

```
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.  
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.226 ms  
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.062 ms  
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.063 ms  
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.104 ms  
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.067 ms  
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.061 ms  
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.063 ms  
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.062 ms  
64 bytes from 192.168.60.5: icmp_seq=25 ttl=63 time=0.060 ms  
64 bytes from 192.168.60.5: icmp_seq=31 ttl=63 time=0.109 ms  
64 bytes from 192.168.60.5: icmp_seq=37 ttl=63 time=0.084 ms  
64 bytes from 192.168.60.5: icmp_seq=42 ttl=63 time=0.199 ms  
64 bytes from 192.168.60.5: icmp_seq=48 ttl=63 time=0.136 ms  
64 bytes from 192.168.60.5: icmp_seq=54 ttl=63 time=0.062 ms  
^C  
--- 192.168.60.5 ping statistics ---  
59 packets transmitted, 14 received, 76.2712% packet loss, time 59500ms
```

去掉第二条规则后则报文数量和原来一样，平均每秒一个，没有受限。因此第二条规则是不可缺的，因为第二条规则决定了如何处理超出第一条限制的报文，没有第二条规则那么路由在处理超出限制的报文时采用接收，从而达不到限制报文数量的作用。

```
--- 192.168.60.5 ping statistics ---  
30 packets transmitted, 30 received, 0% packet loss, time 29891ms  
rtt min/avg/max/mdev = 0.060/0.066/0.083/0.005 ms
```

## Task5: Load Balancing

### Using the nth mode (round-robin)

在 router 上配置 eth 规则后，在 router 和 192.168.60.5 上都开启 nc -luk 8080，不断从 10.9.0.5 向 10.9.0.11 发送 hello，可以看到每三个里有一个发给了 192.168.60.5。

```
root@f71f0b983d48:/# nc -luk 8080  
hello  
hello  
hello  
hello  
hello  
hello  
hello  
root@660e8067c2c5:/# nc -luk 8080  
hello  
hello  
hello  
■
```



## 添加新规则

```
root@f71f0b983d48:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 1 -j DNAT --to-destination 192.168.60.6:8080
root@f71f0b983d48:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 2 -j DNAT --to-destination 192.168.60.7:8080
root@f71f0b983d48:/# nc -luk 8080
```

在这种规则下，数据包按照顺序一人一个 hello

```
root@660e8067c2c5:/# nc root@60979ddce039:/# nc -luk 8080 root@da123ec14e29:/# nc -luk 8080
hello hello hello
hello hello hello
hello hello hello
hello hello hello
hello hello hello
hello hello hello
hello hello hello
hello
```

### Using the random mode.

使用 random 规则，其中向 192.168.60.5 发送报文的概率为 0.5，向 192.168.60.6 发送报文的概率为 0.25，向 192.168.60.7 发送报文的概率为 0.25

```
root@f71f0b983d48:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.5 -j DNAT --to-destination 192.168.60.5:8080
root@f71f0b983d48:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.25 -j DNAT --to-destination 192.168.60.6:8080
root@f71f0b983d48:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.25 -j DNAT --to-destination 192.168.60.7:8080
```

在这种规则下，192.168.60.5、192.168.60.6、192.168.60.7 收到的 hello 数量差不多为 2:1:1。

[illegible]